

# Big Data Architecture

Guido Schmutz



BASEL ■ BERN ■ BRUGG ■ DÜSSELDORF ■ FRANKFURT A.M. ■ FREIBURG I.BR. ■ GENEVA  
HAMBURG ■ COPENHAGEN ■ LAUSANNE ■ MUNICH ■ STUTTGART ■ VIENNA ■ ZURICH

**trivadis**  
makes IT easier. ■ ■ ■

# ■ Guido Schmutz

Working for Trivadis for more than 18 years

Oracle ACE Director for Fusion Middleware and SOA

Co-Author of different books

Consultant, Trainer Software Architect for Java, Oracle, SOA and Big Data / Fast Data

Member of Trivadis Architecture Board

Technology Manager @ Trivadis

More than 25 years of software development experience

Contact: [guido.schmutz@trivadis.com](mailto:guido.schmutz@trivadis.com)

Blog: <http://guidoschmutz.wordpress.com>

Twitter: [gschmutz](https://twitter.com/gschmutz)



**trivadis**  
makes IT easier. ■ ■ ■

# ■ Agenda

1. Introduction
2. Traditional Architecture for Big Data
3. Streaming Analytics Architecture for Fast Data
4. Lambda/Kappa/Unifed Architecture for Big Data
5. Summary

# Introduction

# ■ Big Data is still “work in progress”

Choosing the right architecture is key for any (big data) project

Big Data is still quite a young field and therefore there are no standard architectures available which have been used for years

In the past few years, a few architectures have evolved and have been discussed online

Know the use cases before choosing your architecture

To have one/a few reference architectures can help in choosing the right components

# Hadoop Ecosystem – many choices ....

Unstructured  
Data Sources



Analytics



SQL on  
Hadoop



Management /  
Monitoring



Workflow/Job



Structured Data  
Sources



Core



Data Storage



Serialization



Security



# ■ Important Properties to choose a Big Data Architecture

Latency

Keep raw and un-interpreted data “forever” ?

Volume, Velocity, Variety, Veracity

Ad-Hoc Query Capabilities needed ?

Robustness & Fault Tolerance

Scalability

...

# ■ From Volume and Variety to Velocity

Big Data has evolved ...

## Past

Big Data = Volume & Variety

## Present

Big Data = Volume & Variety & Velocity

and the Hadoop Ecosystem as well ....

## Past

Batch Processing

Time to insight of **Hours**

## Present

Batch & Stream Processing

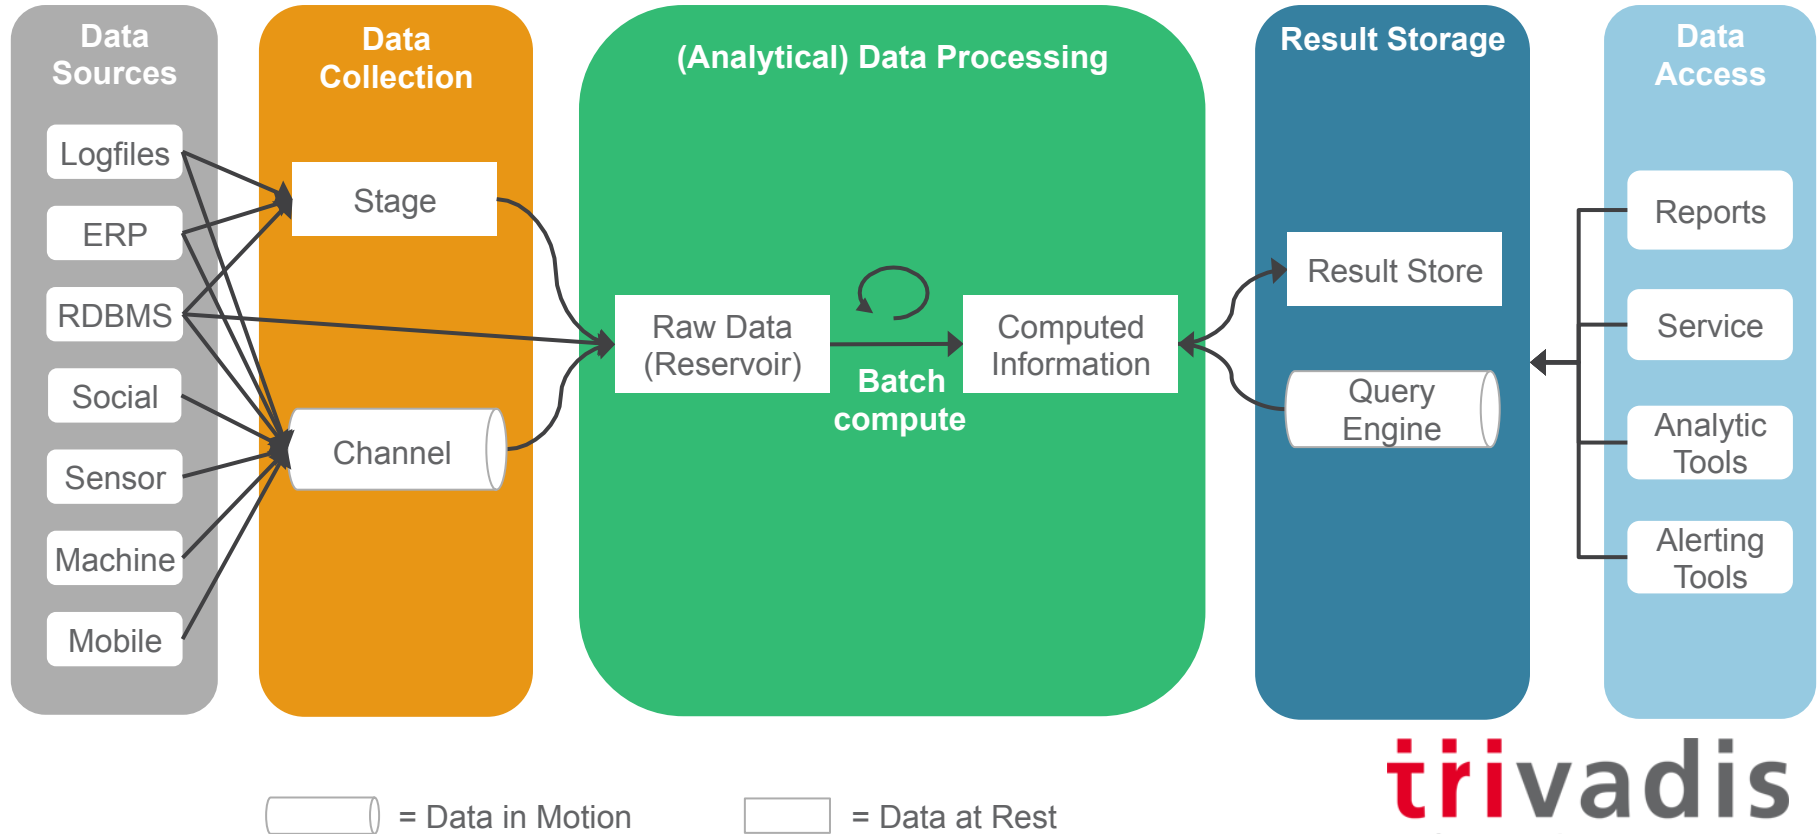
Time to insight in **Seconds**

Adapted from Cloudera Blog article

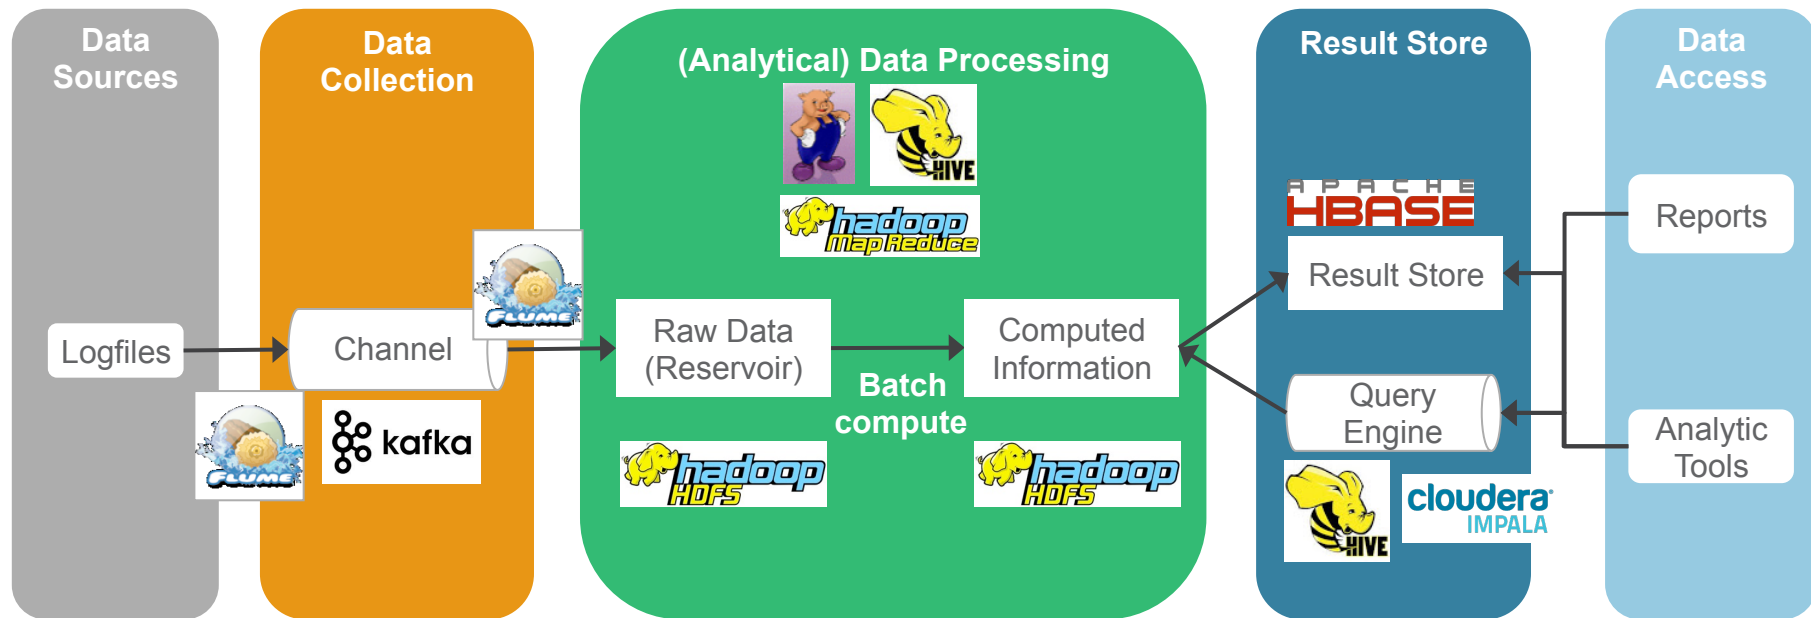


# Traditional Architecture for Big Data

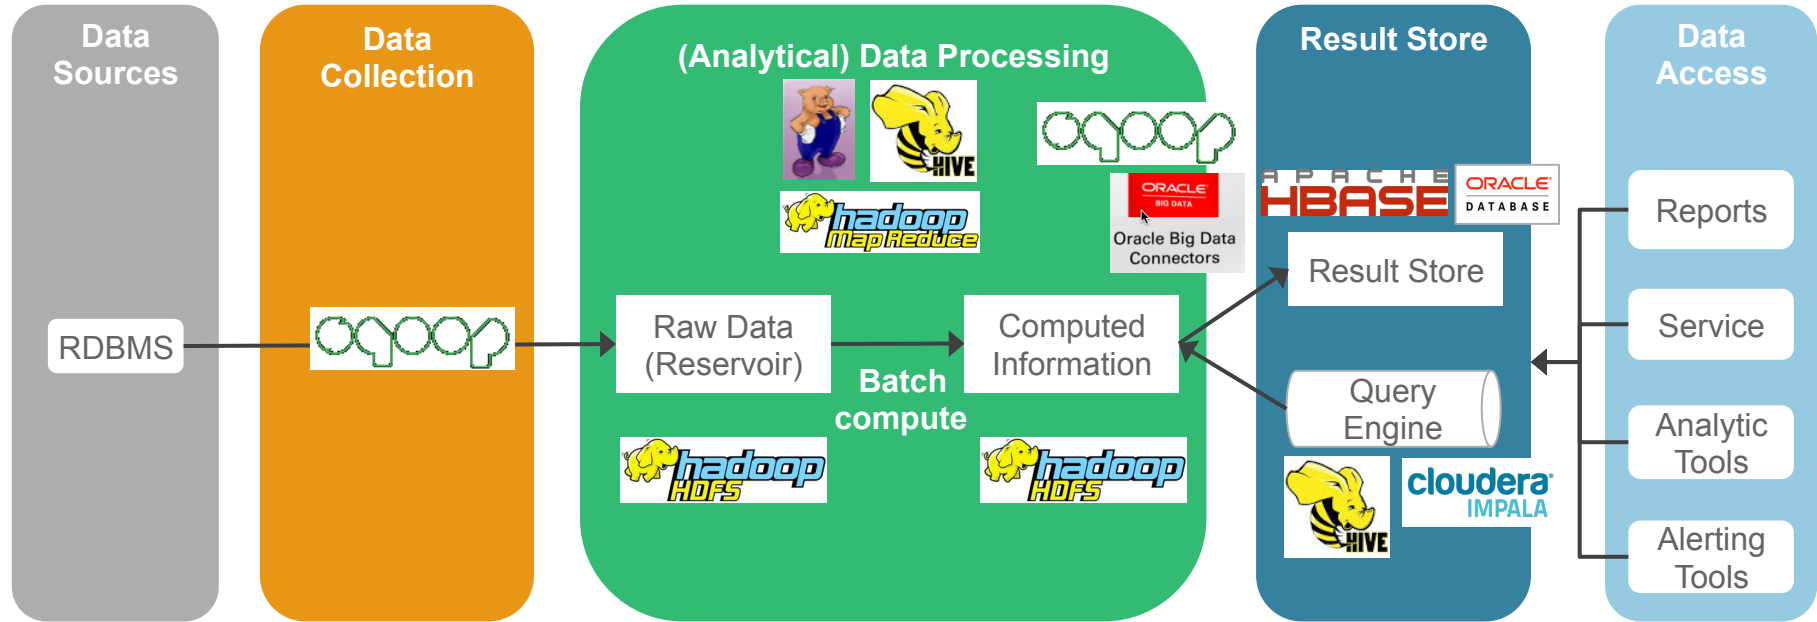
# ■ “Traditional Architecture” for Big Data



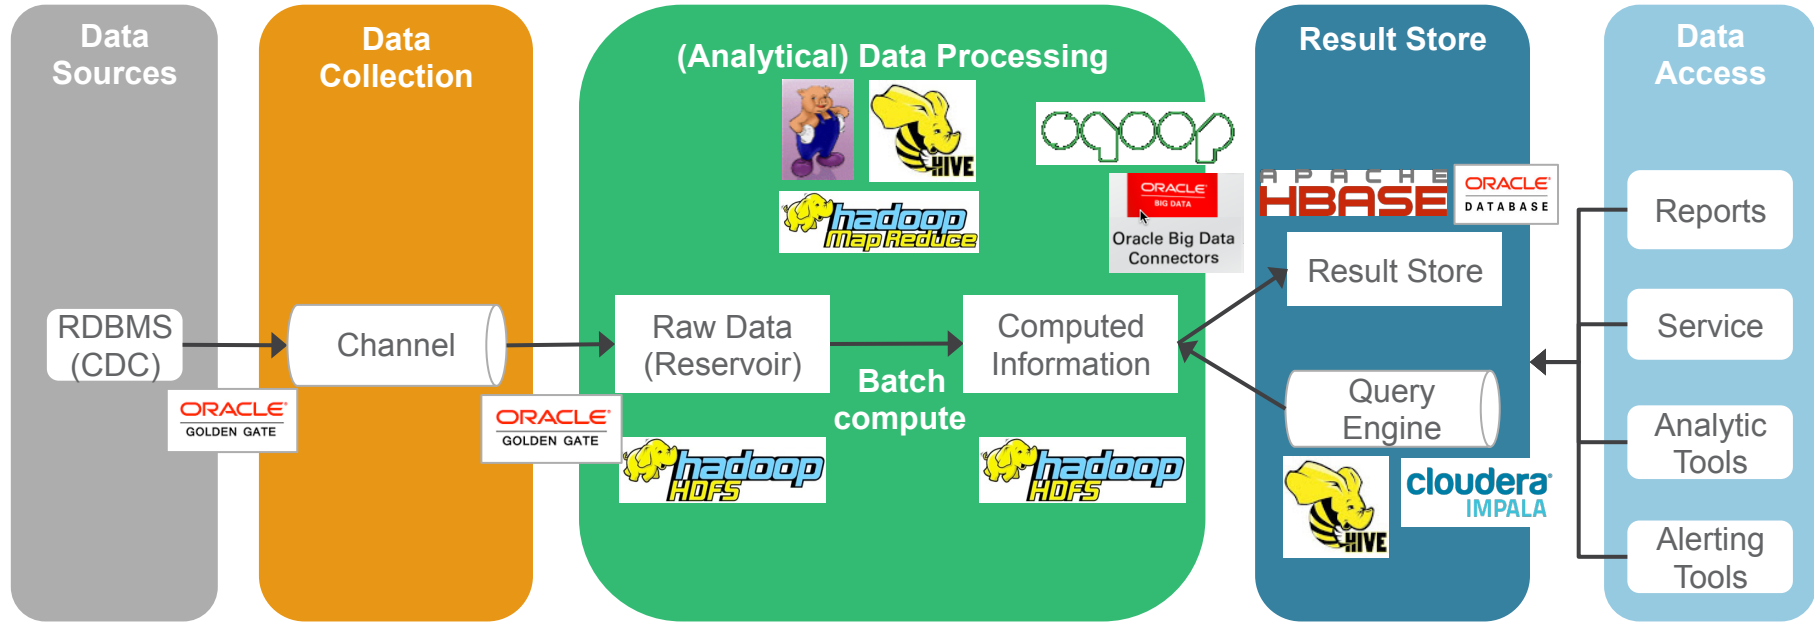
# ■ Use Case 1) – Click Stream analysis: 360 degree view of customer



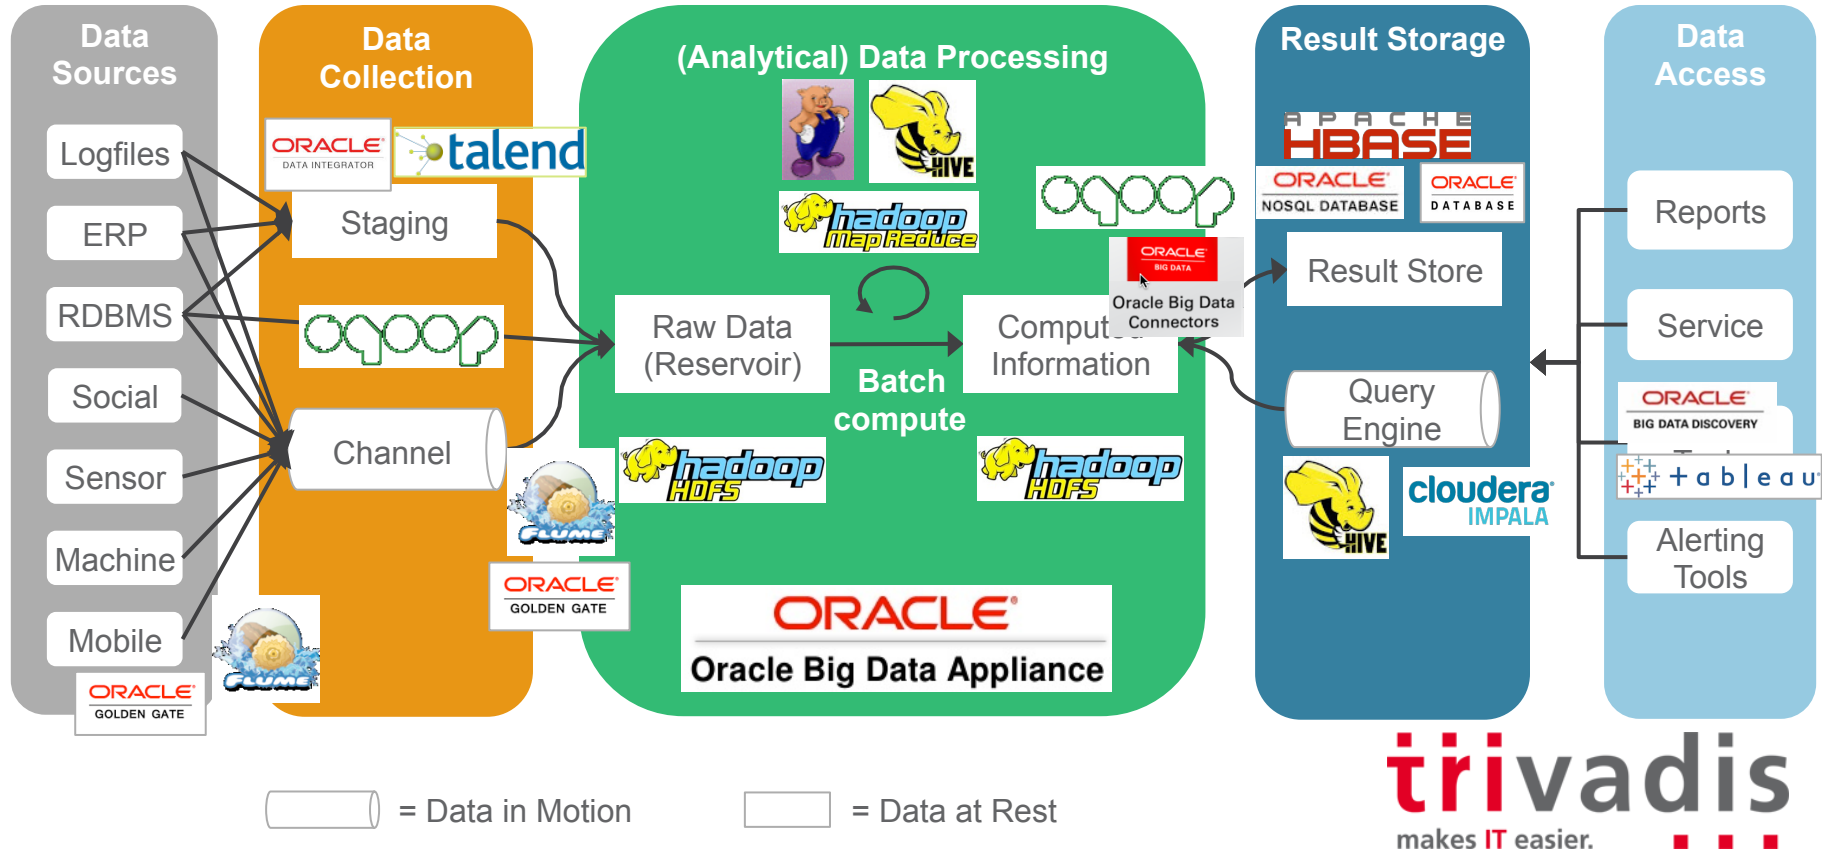
## ■ Use Case 2) – Ingest Relational Data into Hadoop and make it accessible



## ■ Use Case 2a) – Ingest Relational Data into Hadoop and make it accessible




# “Hadoop Ecosystem” Technology Mapping



# ■ Apache Spark – the new kid on the block



Apache Spark is a fast and general engine for large-scale data processing


- The hot trend in Big Data!
- Originally developed 2009 in UC Berkley's AMPLab The AMPLab logo, consisting of the word "amplab" in a dark blue, sans-serif font, followed by a stylized orange line graph icon.
- Can run programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk
- One of the largest OSS communities in big data with over 200 contributors in 50+ organizations
- Open Sourced in 2010 – since 2014 part of Apache Software foundation
- Supported by many vendors

# ■ Motivation – Why Apache Spark?

Apache Hadoop MapReduce: Data Sharing on Disk



Apache Spark: Speed up processing by using Memory instead of Disks





# ■ Apache Spark “Ecosystem”

## Libraries

Spark SQL  
(Batch  
Processing)

Blink DB  
(Approximate  
Querying)

Spark Streaming  
(Real-Time)

MLlib, Spark R  
(Machine  
Learning)

GraphX  
(Graph  
Processing)

## Core Runtime

Spark Core API and Execution Model

## Cluster Resource Managers

Spark  
Standalone

MESOS

YARN

## Data / Data Stores

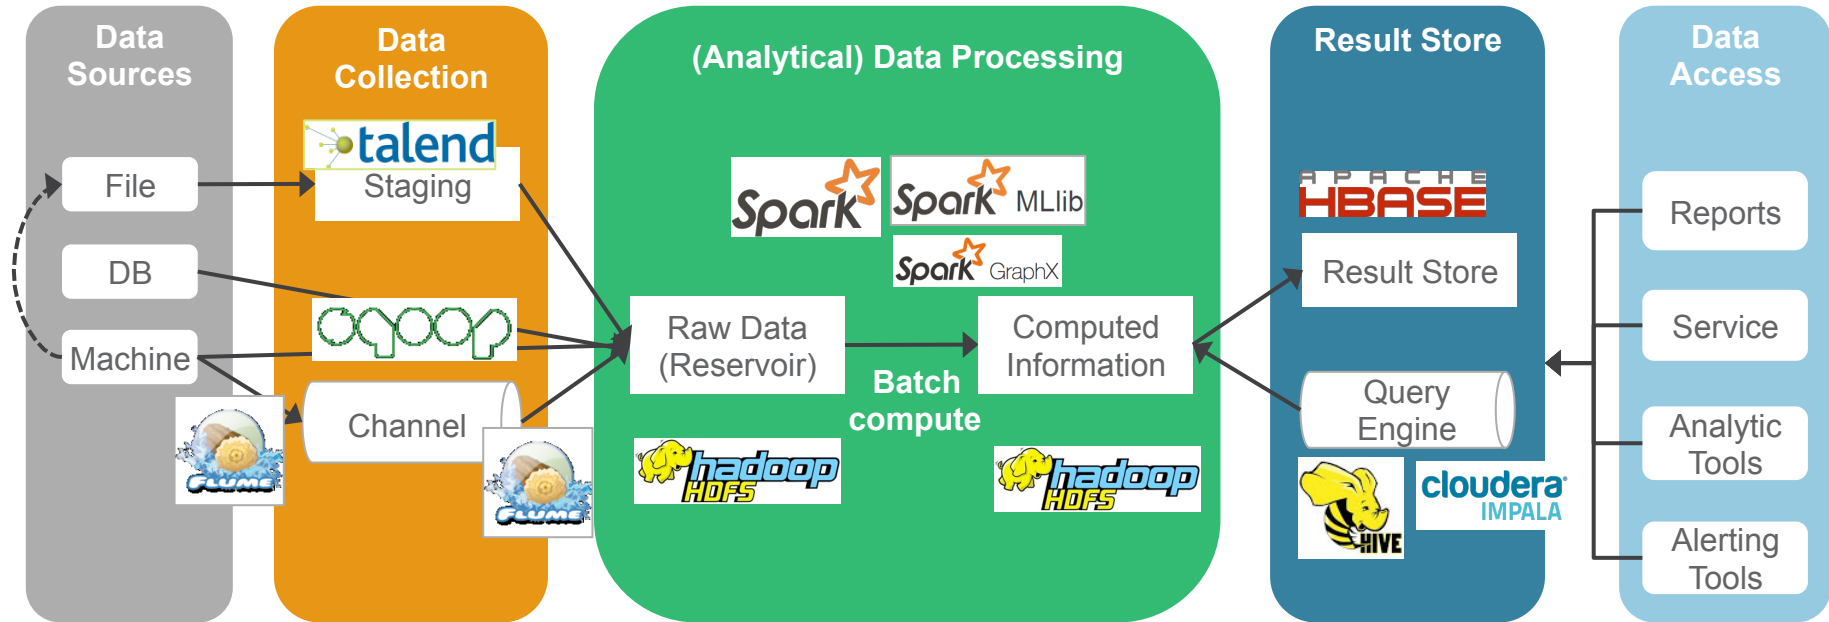
HDFS

Elastic  
Search

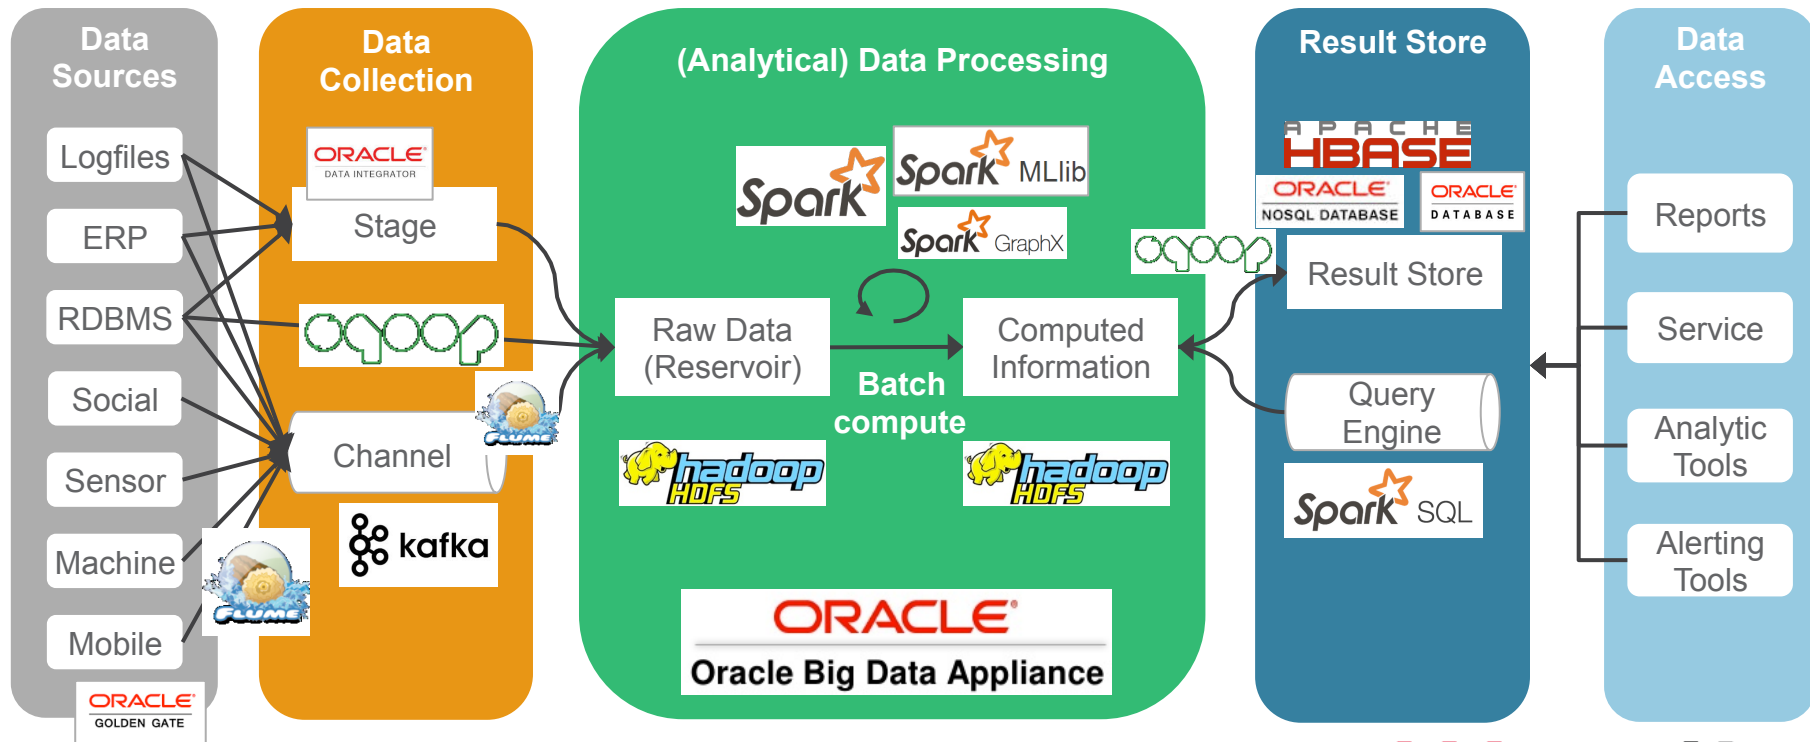
NoSQL

S3

## ■ Use Case 3) – Predictive Maintenance through Machine Learning on collected data



# “Spark Ecosystem” Technology Mapping



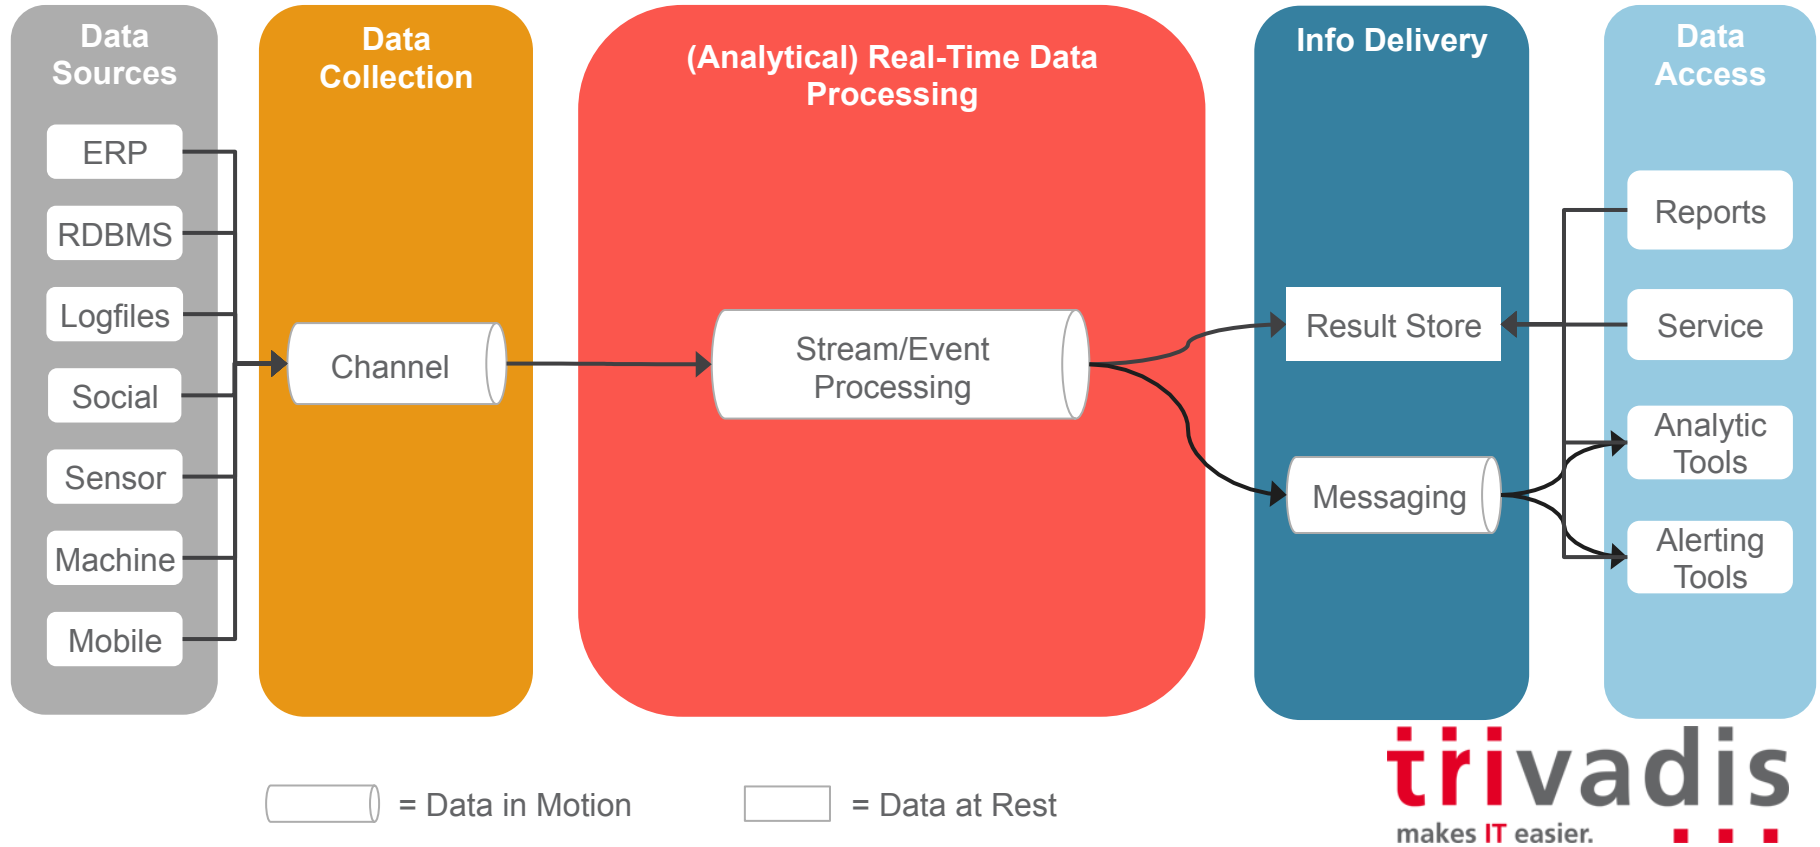
# ■ Traditional Architecture for Big Data

- Batch Processing
- Not for low latency use cases
- Spark can speed up, but if positioned as alternative to Hadoop Map/Reduce, it's still Batch Processing
- Spark Ecosystems offers a lot of additional advanced analytic capabilities (machine learning, graph processing, ...)

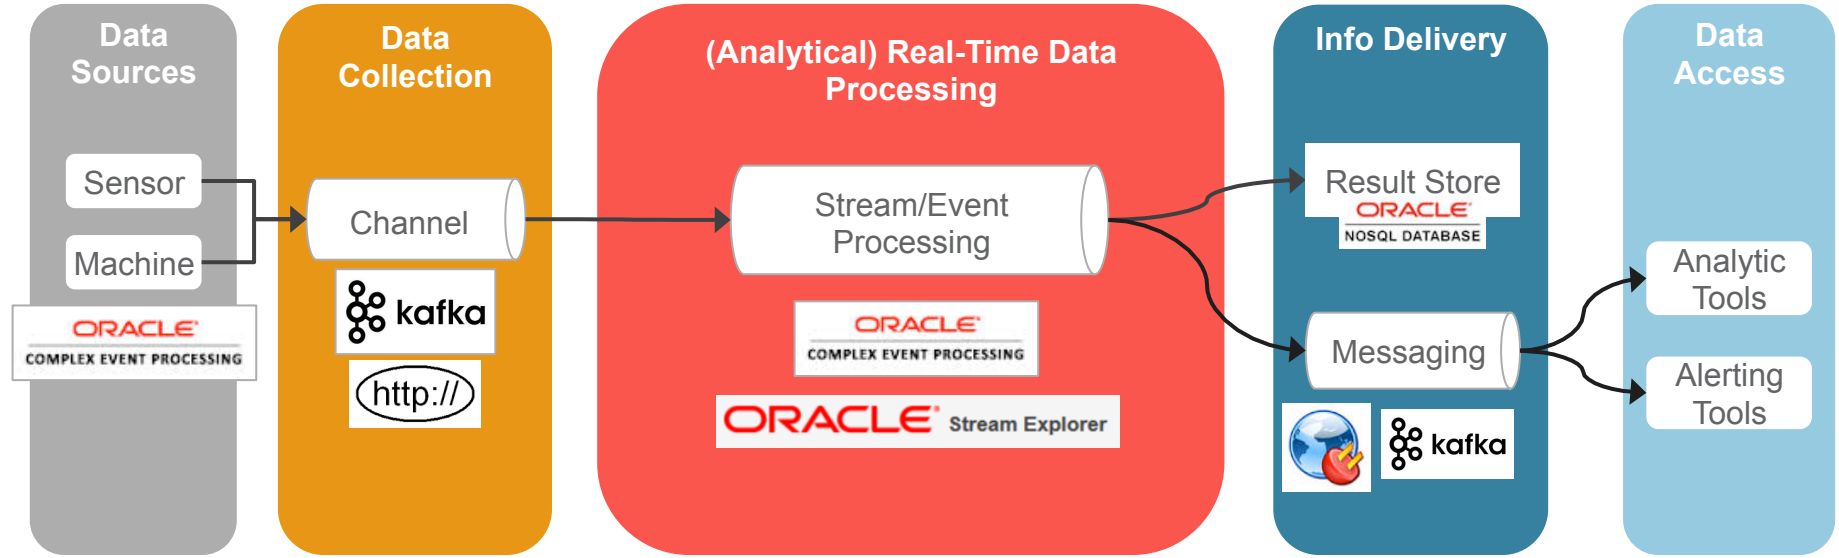
# Streaming Analytics Architecture for Big Data

# Streaming Analytics Architecture for Big Data

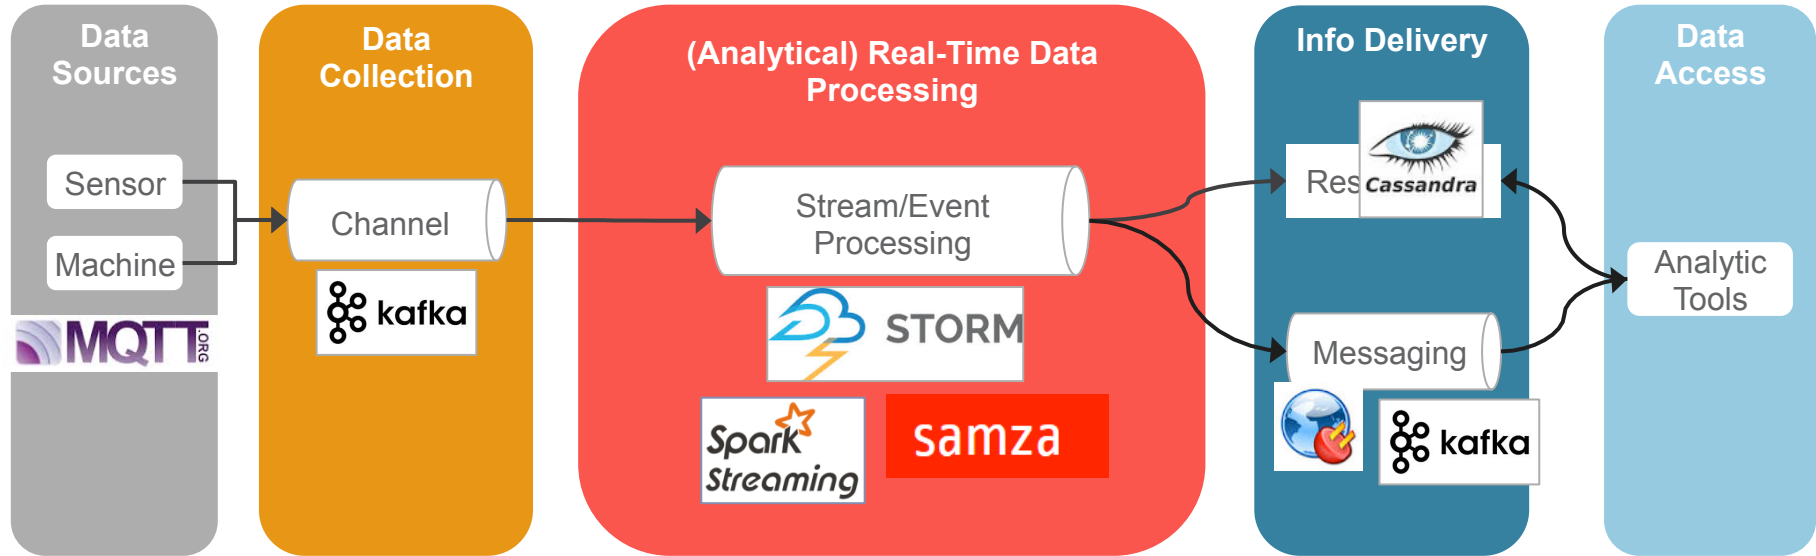
aka. (Complex) Event Processing



## ■ Use Case 4) Alerting in Internet of Things (IoT)



## ■ Use Case 5) Real-Time Analytics on Sensor Events



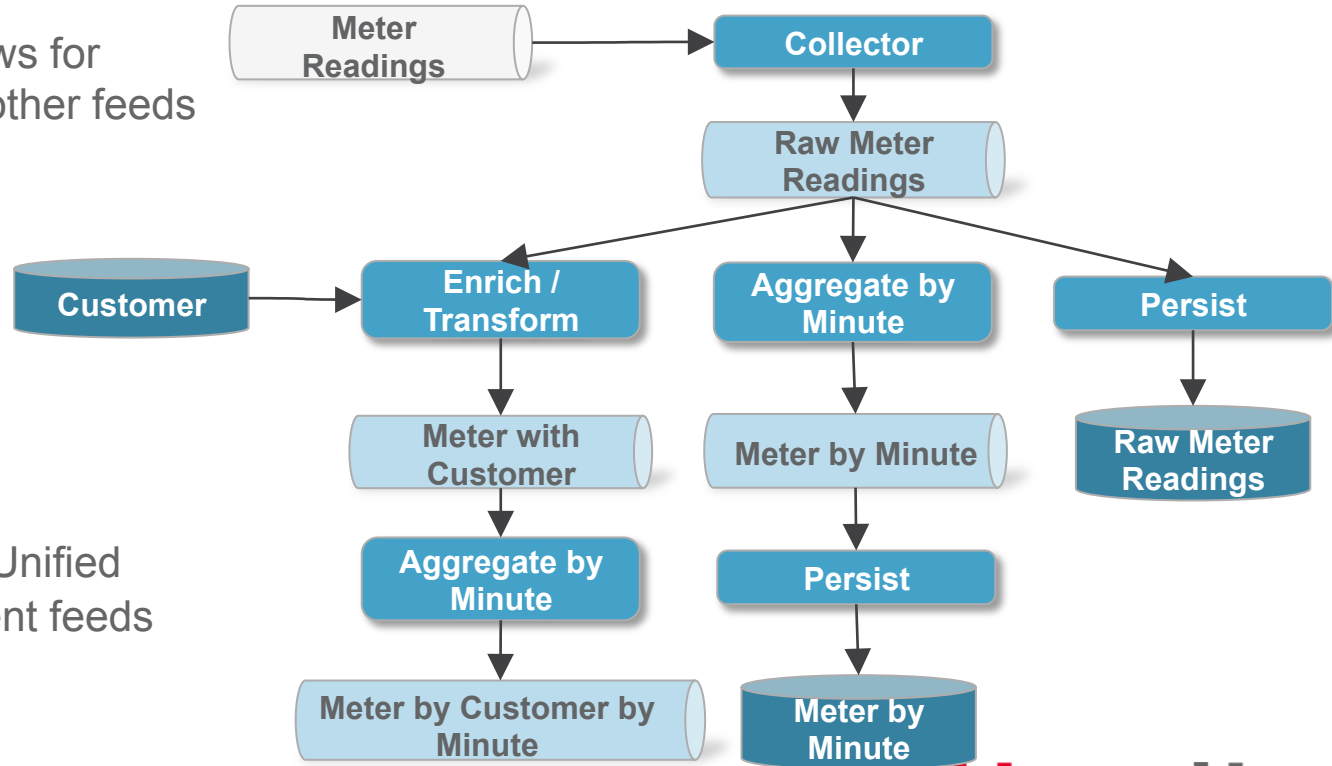


# Unified Log (Event) Processing

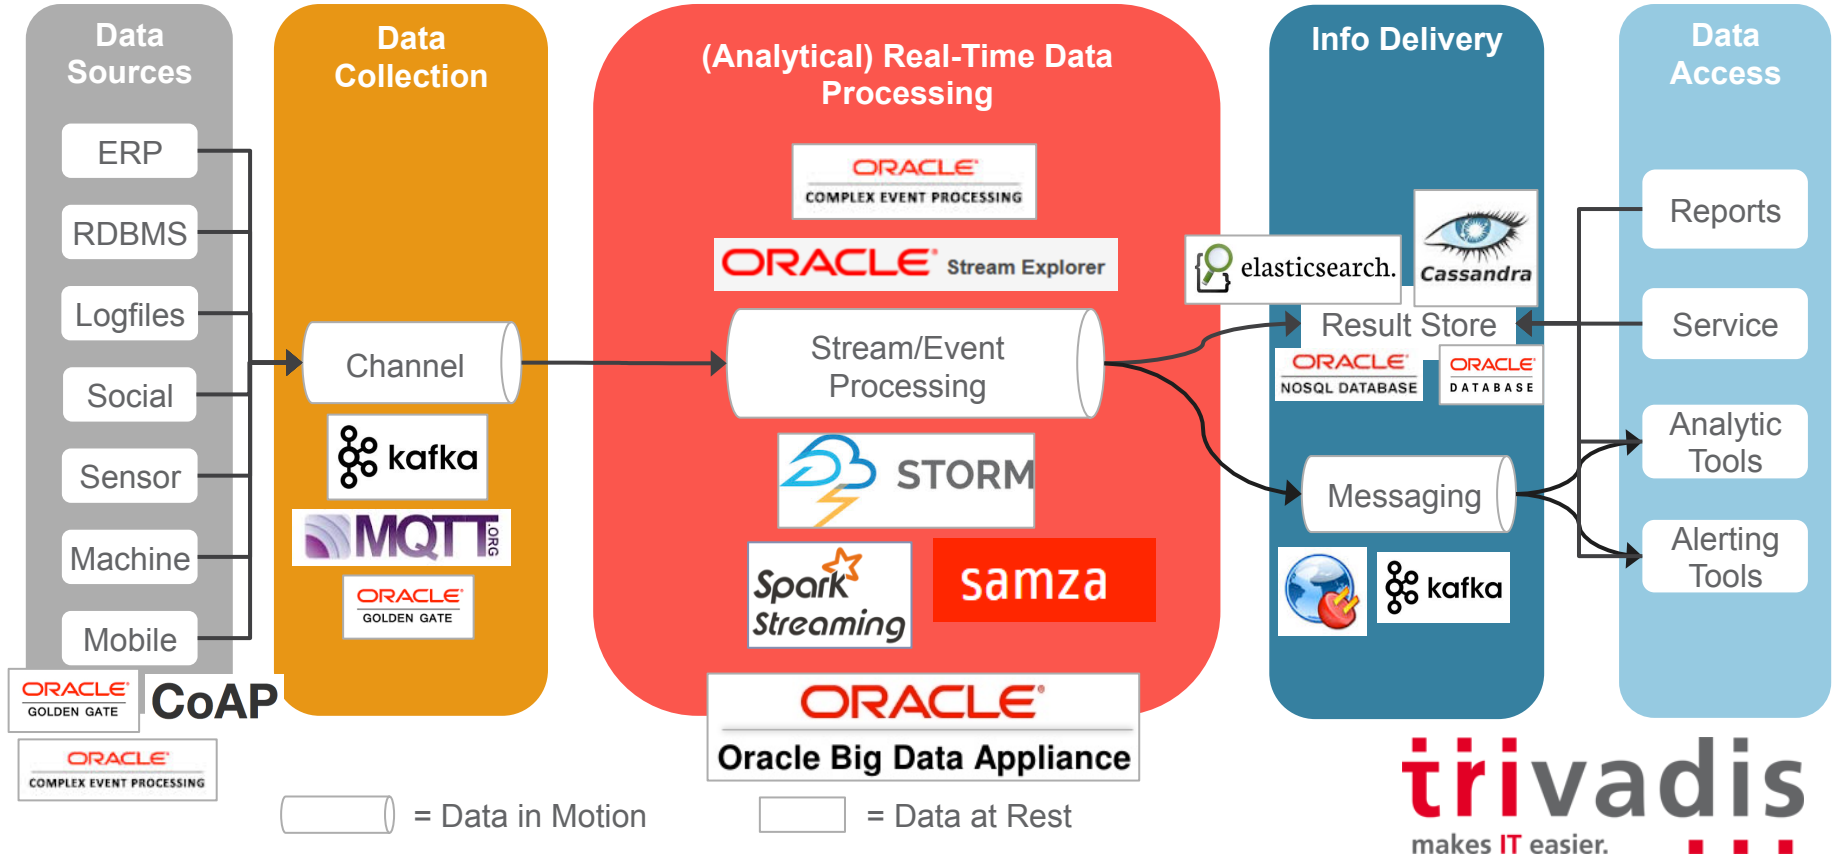
Stream processing allows for computing feeds off of other feeds

Derived feeds are no different than original feeds they are computed off

Single deployment of “Unified Log” but logically different feeds



# Streaming Analytics Technology Mapping



# ■ Streaming Analytics Architecture for Big Data

The solution for low latency use cases

Process each event separately => low latency

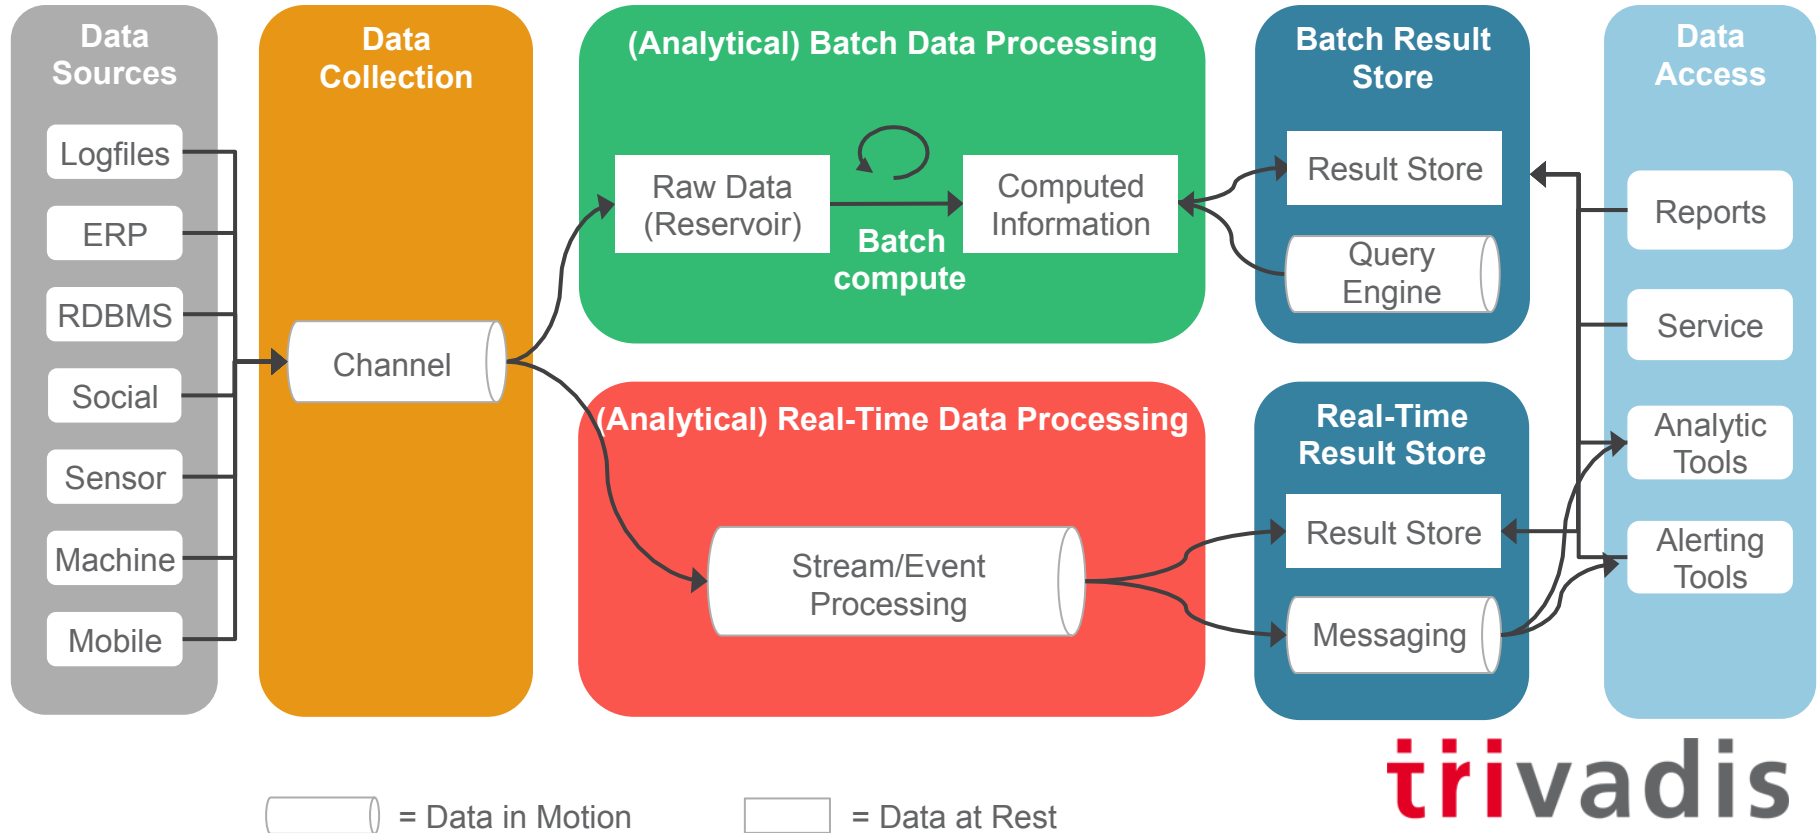
Process events in micro-batches => increases latency but offers better reliability

Previously known as “Complex Event Processing”

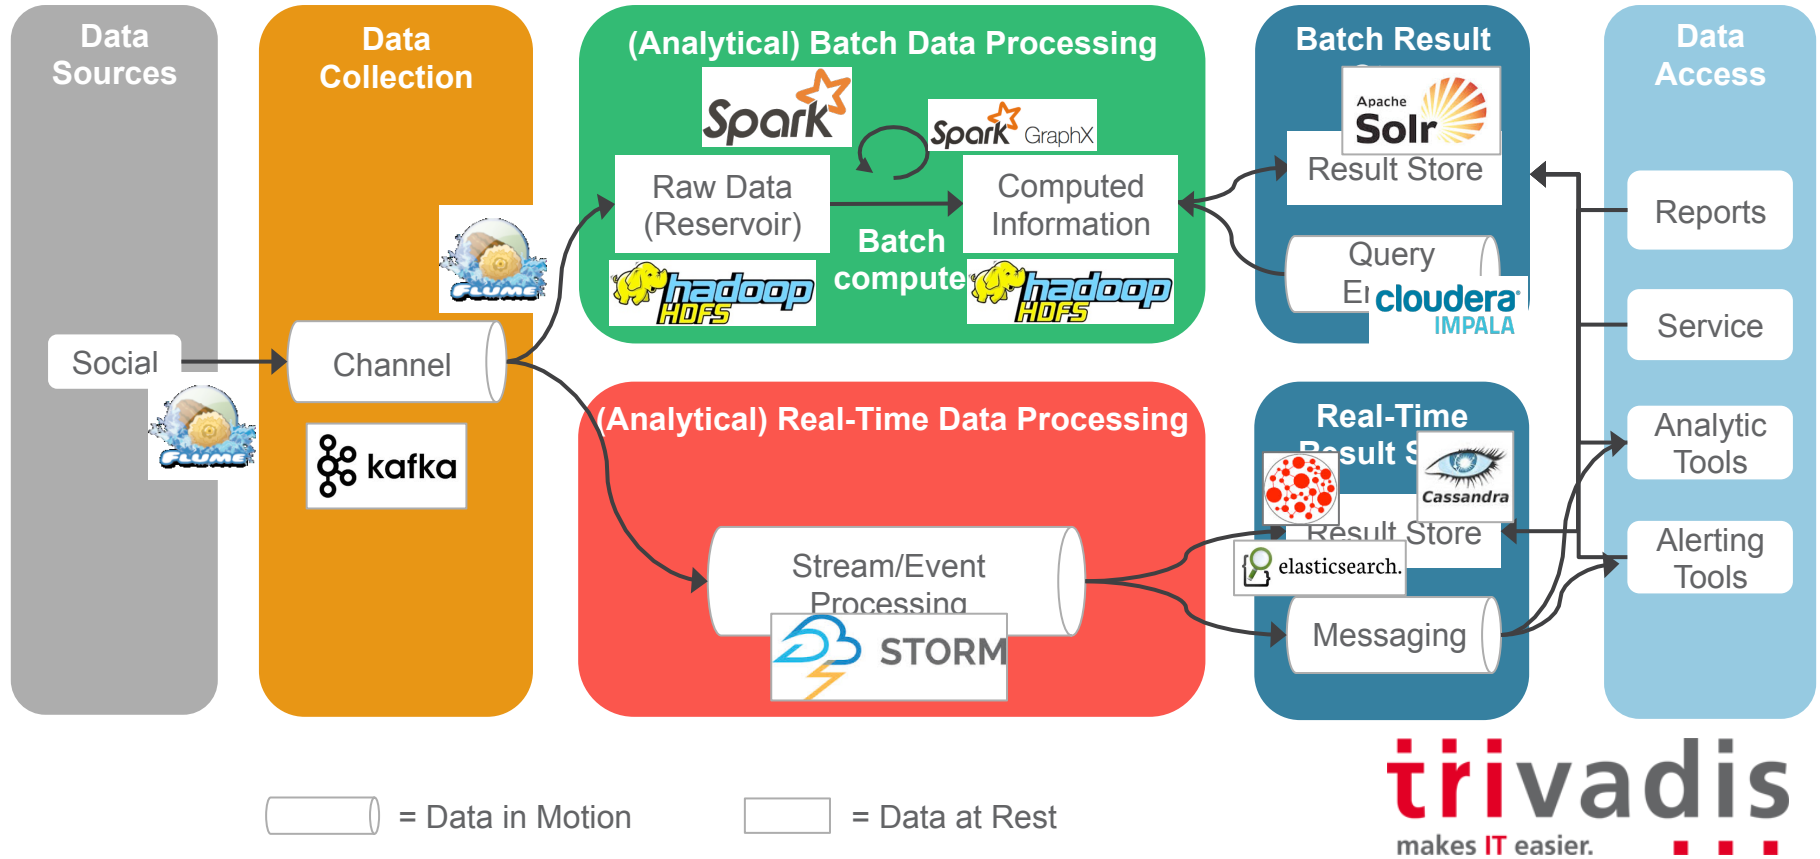
Keep the data moving / Data in Motion instead of Data at Rest => raw events are (often) not stored

# Lambda Architecture for Big Data

# ■ “Lambda Architecture” for Big Data



## ■ Use Case 6) Social Media and Social Network Analysis



# ■ Lambda Architecture for Big Data

Combines (Big) Data at Rest with (Fast) Data in Motion

Closes the gap from high-latency batch processing

Keeps the raw information forever

Makes it possible to rerun analytics operations on whole data set if necessary

=> because the old run had an error or

=> because we have found a better algorithm we want to apply

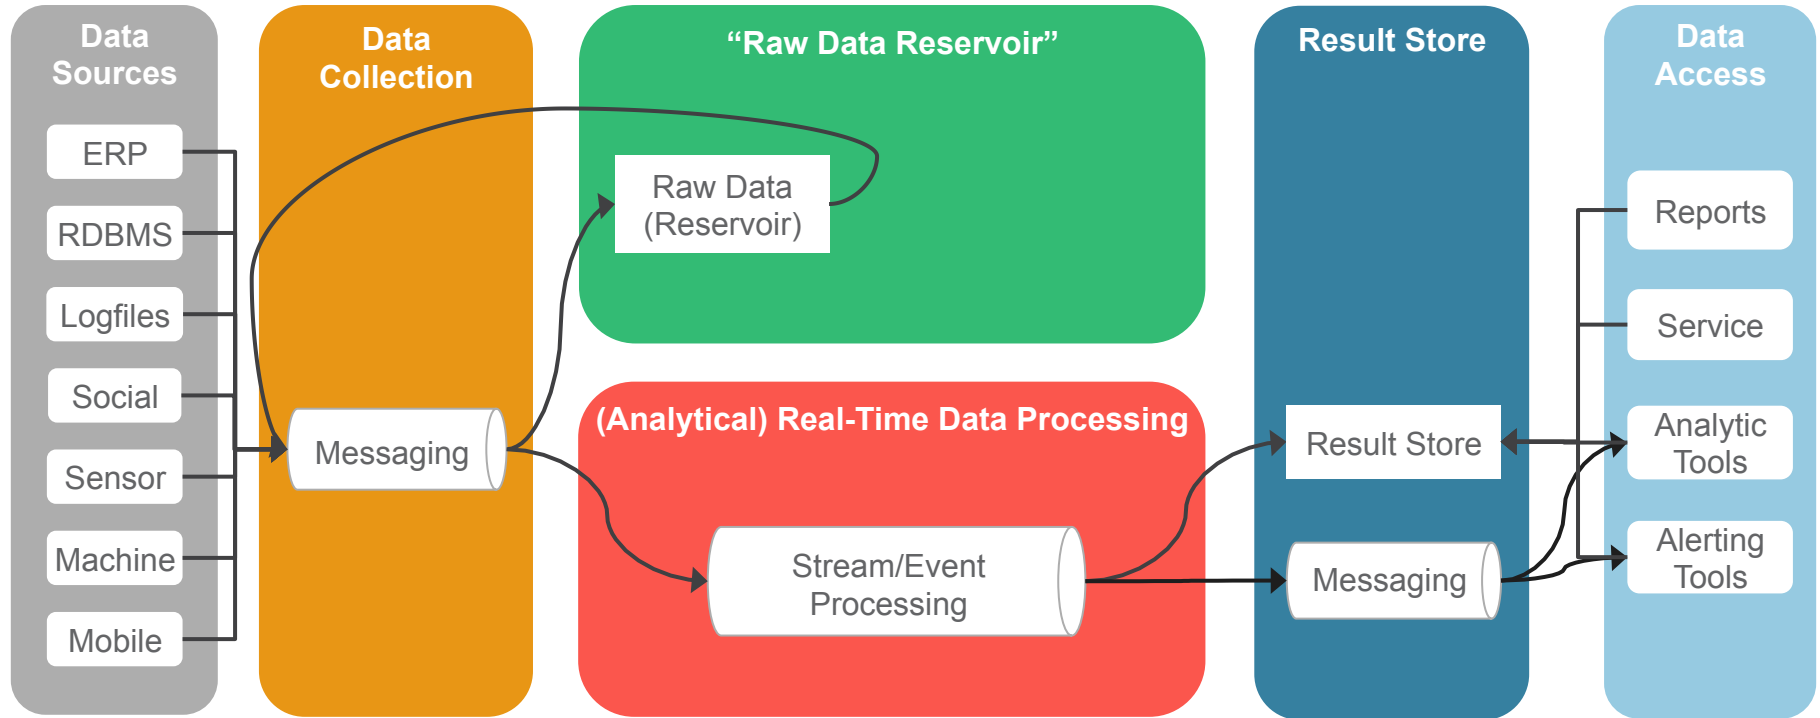
Have to implement functionality twice

- Once for batch
- Once for real-time streaming

# „Kappa“ Architecture for Big Data



# ■ “Kappa Architecture” for Big Data



# ■ Kappa Architecture for Big Data

The solution for low latency use cases

Process each event separately => low latency

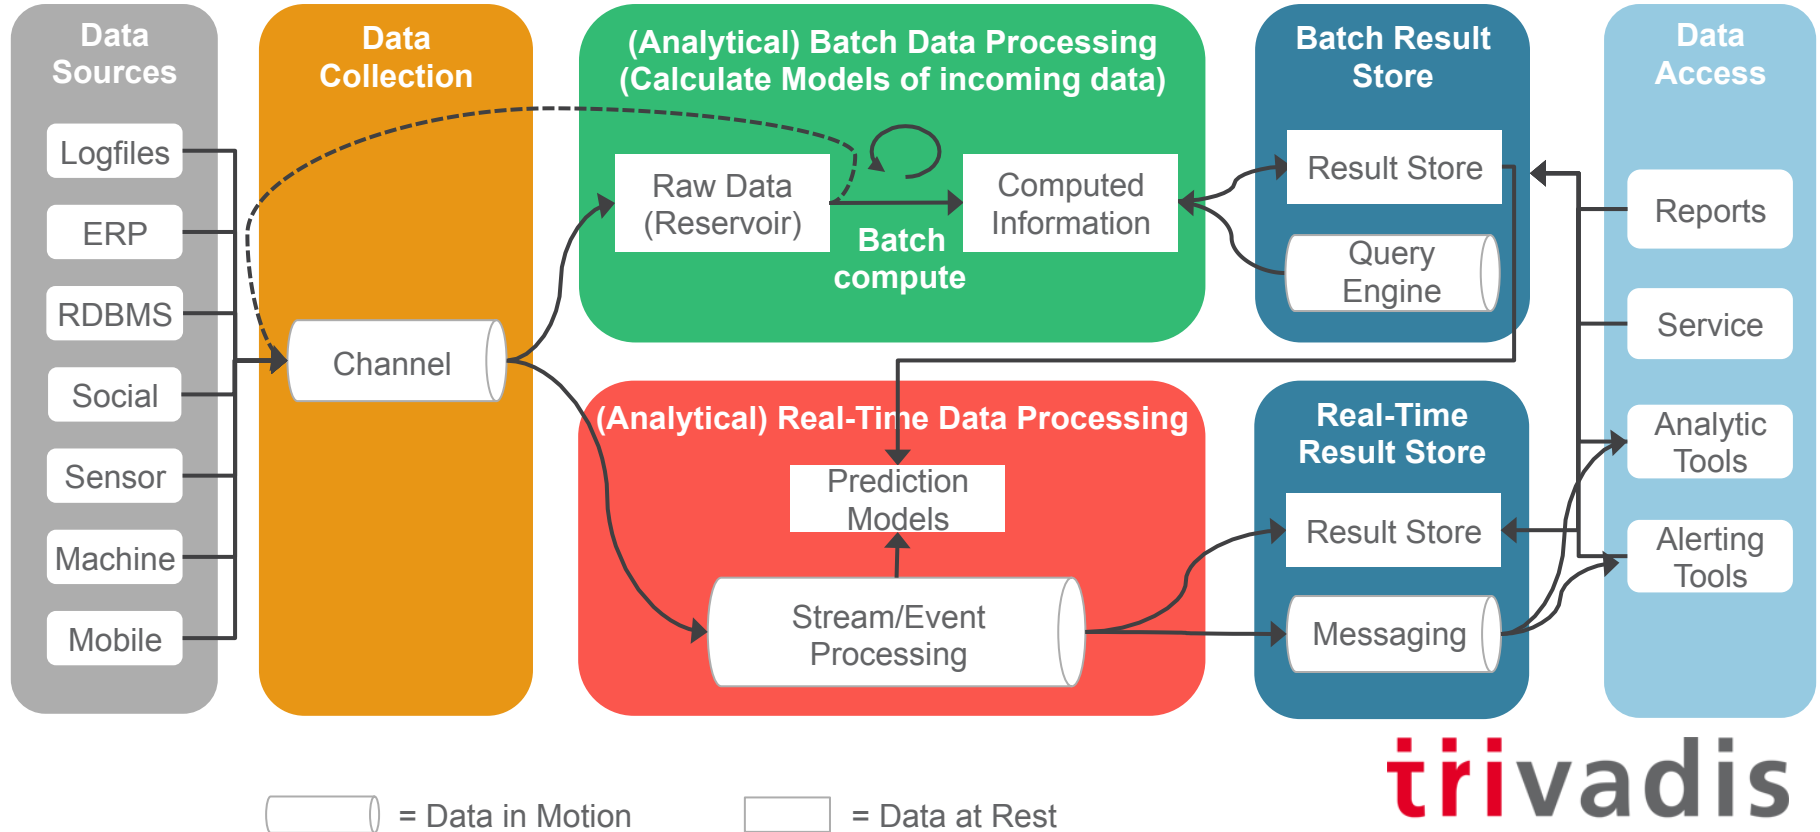
Process events in micro-batches => increases latency but offers better reliability

Previously known as “Complex Event Processing”

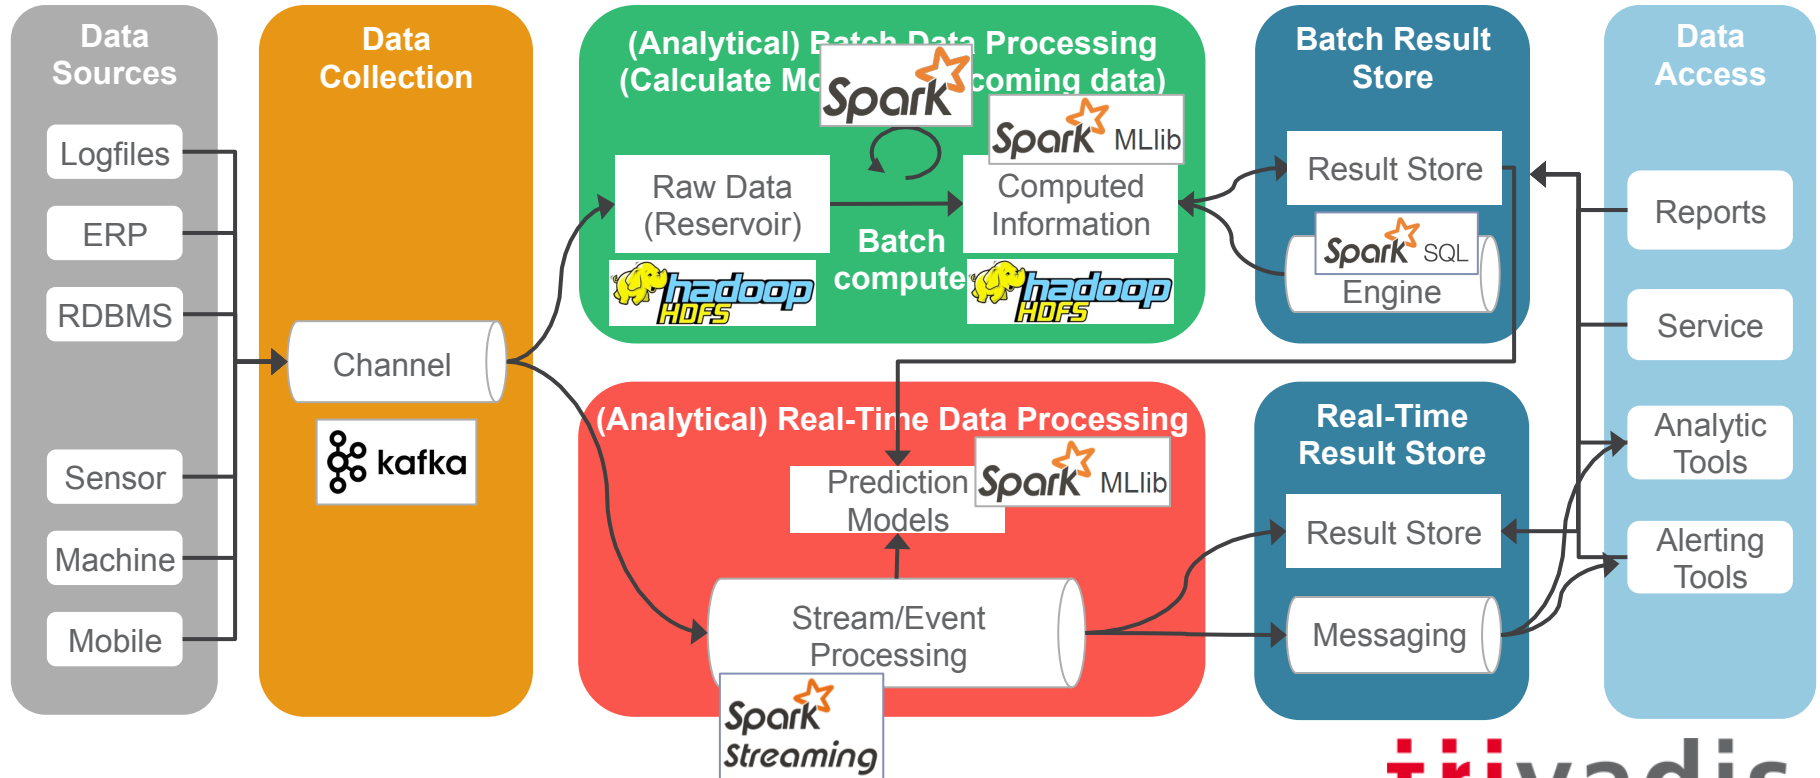
Keep the data moving / Data in Motion instead of Data at Rest

# „Unified“ Architecture for Big Data

# ■ “Unified Architecture” for Big Data



## ■ Use Case 7) Fraud Detection



# Summary

## ■ Summary

Know your use cases and then choose your architecture and the relevant components/products/frameworks

You don't have to use all the components of the Hadoop Ecosystem to be successful

Big Data is still quite a young field and therefore there are no standard architectures available which have been used for years

Lambda, Kappa Architecture are best practices architectures which you have to adapt to your environment



**trivadis**  
makes **IT** easier. ■ ■ ■



Name Referent  
Titel Referent

Tel. +00 00 000 00 00

vorname.name@trivadis.com

